

Question 4 (18 points):

When generating initial code, a compiler generates naive code where the instructions for the execution of each memory access or for the computation of each expression is done independently of other such accesses or computation in the code. Consider the following function written in C to compute the Hadamard product of two vectors:

```
void Hadamard(int *A, int *B, int N){
    for(i=0; i<N; i++){
        A[i] = A[i] * B[i];
    }
}
```

<pre>7 HadamardNaive: 8 add t0, zero, zero # i <- 0 9 loop: bge t0, a2, done # if i >= N goto done 10 slli t1, t0, 2 # t1 <- 4*i 11 add t2, a0, t1 # t2 <- Address(A[i]) 12 lw t3, 0(t2) # t3 <- A[i] 13 slli t4, t0, 2 # t4 <- 4*i 14 add t5, a1, t4 # t5 <- Address(B[i]) 15 lw t6, 0(t5) # t6 <- B[i] 16 mul t7, t3, t6 # t7 <- A[i]*B[i] 17 slli t8, t0, 2 # t8 <- 4*i 18 add t9, a0, t8 # t9 <- Address(B[i]) 19 sw t7, 0(t9) # A[i] <- A[i]*B[i] 20 addi t0, t0, 1 # i <- i+1 21 jal zero, loop 22 done: jalr zero, ra, 0</pre>	<pre>24 HadamardOpt: 25 slli t0, a2, 2 # t0 <- 4*N 26 add t1, a0, t0 # t1 <- Address(A[N]) 27 bge zero, a2, done # if(0<=N) done 28 loop: lw t2, 0(a0) # t2 <- A[i] 29 lw t3, 0(a1) # t3 <- B[i] 30 mul t4, t2, t3 # t4 <- A[i]*B[i] 31 sw t4, 0(a0) # A[i] <- A[i]*B[i] 32 addi a0, a0, 4 # pA++ 33 addi a1, a1, 4 # pB++ 34 blt a0, t1, loop # if pA < Address(A[N]) 35 done: jalr zero, ra, 0</pre>
(a) Naive Version	(b) Optimized Version

Figure 1: Two versions for a Dot Product function.

Figure ?? shows the RISC-V assembly code for a naive version and for an optimized version of the this function. In this question you will study the performance of these two versions. A performance study was conducted to determine the average number of clock cycles for different types of instructions. Based on this study, the instructions used in Figure ?? are classified into the following classes of instructions:

ALU Instructions (add, slli, addi): 1 cycle

Jumps and Branches (bge, blt, jal, jalr): 3 cycles

Memory instructions (lw, sw): 5 cycles

Multiplication Instructions (mul): 10 cycles

- a. (3 points) Assuming that N is very large, what is the CPI of the naive version and what is the CPI of the optimized version?

- b. (4 points) Again, assuming that N is very large, which version is faster and by how much?
- c. (5 points) A given invocation of the optimized version of **Hadamard** executes in 15 seconds in a baseline machine with a clock cycle of 2 GHz ($1 \text{ GHz} = 10^9 \text{ Hz}$). What was the value of N for this invocation of the **Hadamard** function?
- d. (6 points) A new version of the same processor has been designed that implements the following changes:
- The clock frequency is 3 GHz
 - The average number of cycles required to execute memory instructions is reduced to 3 cycles
 - The number of cycles required to execute a multiplication instruction is also reduced

To discover what is the average number of cycles required to execute a multiplication operation in this new version of the machine you run an experiment where you execute the optimized version of **Hadamard** with $N = 2 \times 10^6$. You find out that the execution of such an invocation of **Hadamard** takes 12.7 ms ($1 \text{ ms} = 10^{-3} \text{ seconds}$). What is the average number of cycles required for multiplication operations in this new version of the machine?